



AALBORG
UNIVERSITY

Database System (SW5)

2. Relational Model

Tiantian Liu
Department of Computer Science
Aalborg University
Fall 2025



Motivation

- SQL's data model resembles the relational model
- SQL is partly based on relational algebra
- Relational algebra helps understanding query plans and query optimization



Learning Goals

- Explain the relational model
- Create non-trivial relational algebra queries
- Use the various join types in relational algebra queries

Agenda

- Relational Model
 - Relation, tuple, attribute
 - Schema and instance
 - Keys: Superkey, candidate key, primary key, foreign key
- Relational Algebra



Relation

- In the relational model the term **relation** is used to refer to a table

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

- The *Instructor Relation*

Tuple

- The term **tuple** is used to refer to a row.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

- The *Instructor Relation*

Relations are Unordered

- Order of tuples is irrelevant (tuples may be stored in an arbitrary order)

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000	22222	Einstein	Physics	95000
12121	Wu	Finance	90000	12121	Wu	Finance	90000
15151	Mozart	Music	40000	32343	El Said	History	60000
22222	Einstein	Physics	95000	45565	Katz	Comp. Sci.	75000
32343	El Said	History	60000	98345	Kim	Elec. Eng.	80000
33456	Gold	Physics	87000	76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000	10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000	58583	Califieri	History	62000
76543	Singh	Finance	80000	83821	Brandt	Comp. Sci.	92000
76766	Crick	Biology	72000	15151	Mozart	Music	40000
83821	Brandt	Comp. Sci.	92000	33456	Gold	Physics	87000
98345	Kim	Elec. Eng.	80000	76543	Singh	Finance	80000

Attribute

- The term **attribute** refers to a column of a table

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

- The *Instructor* Relation
- Attribute: *ID, name, dept_name, salary*

Attribute--Domain

- For each attribute of a relation, there is a set of permitted values, called the **domain** of that attribute.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

- The *Instructor* Relation
- The domain of the *dept_name* attribute is the set of all possible department names.



Atomic Domain

- Why should a domain be atomic?

ID	Name	Address
001	Elodie	59, SQL Street
002	Felix	103, HTML Road
003	Soren	32, PHP Street

- Domain of *Address* attribute
 - (59, SQL Street), (103, HTML Road), (32, PHP Street)



Atomic Domain

- Better Table

ID	Name	Street_num	Street_name
001	Elodie	59	SQL Street
002	Felix	103	HTML Road
003	Soren	32	PHP Street

- When the data is atomic, that means the data has been broken down into the smallest pieces that can't or shouldn't be divided.
 - Not necessarily smallest possible.
- Again, this depends on your requirement on getting information.
 - Being atomic in one scenario may not be so in another.

Relation Schema and Instance

- A_1, A_2, \dots, A_n are **attributes**
- $R = (A_1, A_2, \dots, A_n)$ is a **relation schema**
 Example
 $\text{instructor} = (ID, name, dept_name, salary)$
- A **relation instance** r defined over schema R is denoted by $r(R)$.
- An element t of relation r is called a **tuple** and is represented by a row in a table.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Relation instance

Database Schema

- Database schema -- is the logical structure of the database.
- Database instance -- is a snapshot of the data in the database at a given instant in time.
- Example: University Database
 - *instructor (ID, name, dept_name, salary)*
 - *course (course_id, title, dept_name ,credits)*
 - *department (dept_name, building, budget)*
 - *section (course_id, sec_id, semester, year, building, room_number, time_slot_id)*
 - *teaches (ID, course_id, sec_id, semester, year)*
 - *student (ID, name, dept_name, tot_cred)*
 - *advisor (s_id, i_id)*
 - *takes (ID, course_id, sec_id, semester, year, grade)*
 - *classroom (building, room_number, capacity)*
 - *time slot (time_slot_id, day, start_time, end_time)*

Superkey

- Let $K \subseteq R$
- K is a superkey of R if values for K are sufficient to identify a unique tuple of each possible relation $r(R)$
- If t_1 and t_2 are in r and $t_1 \neq t_2$, then $t_1.K \neq t_2.K$.
- Example: *instructor* (ID , $name$, $dept_name$, $salary$)
 - $\{ID\}$
 - $\{name\}$
 - $\{dept_name\}$
 - $\{salary\}$
 - $\{ID, name\}$
 - $\{ID, name, dept_name\}$
 - $\{dept_name, salary\}$
 - $\{ID, name, dept_name, salary\}$

Candidate key

- Superkey K is a **candidate key** if K is minimal
 - A superkey may contain extraneous attributes.
 - We are often interested in superkeys for which no proper subset is a superkey.
- Example: *instructor* (ID , $name$, $dept_name$, $salary$)
 - $\{ID\}$, $\{ID, name\}$, $\{ID, name, dept_name\}$ are superkey
 - $\{ID\}$ is a candidate key for *Instructor*
 - $\{ID, name\}$, $\{ID, name, dept_name\}$ are not candidate keys
- Suppose that a combination of $name$ and $dept_name$ is sufficient to distinguish among members of the *instructor* relation
 - Both $\{ID\}$ and $\{name, dept_name\}$ are candidate keys

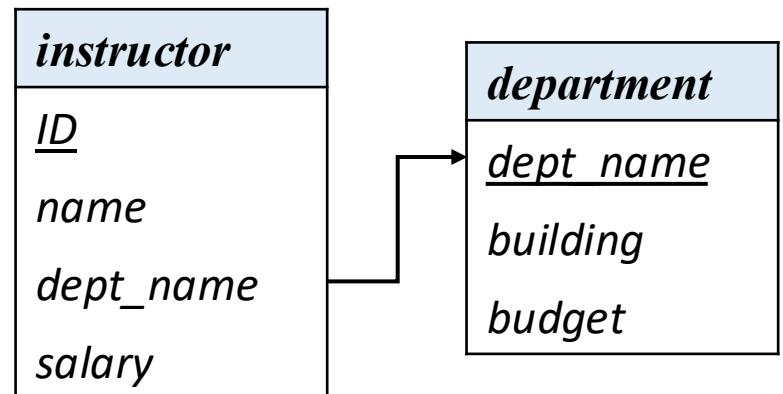


Primary key

- One of the candidate keys is selected to be the primary key.
- Which one?
 - Primary keys must be chosen with care
 - The primary key should be chosen such that its attribute values are never, or are very rarely, changed
- Examples
 - *instructor* (ID, name, dept_name, salary)
 - *course* (course_id, title, dept_name, credits)
 - *department* (dept_name, building, budget)
 - *student* (ID, name, dept_name, tot_cred)
 - *classroom* (building, room_number, capacity)

Foreign key

- Foreign key constraint: Value in one relation must appear in another
- Example:
 - *instructor* (*ID*, *name*, *dept_name*, *salary*)
 - *department* (*dept_name*, *building*, *budget*)
 - *dept_name* in *instructor* is a foreign key from *instructor* referencing *department*
- Referencing relation: *instructor*
- Referenced relation: *department*



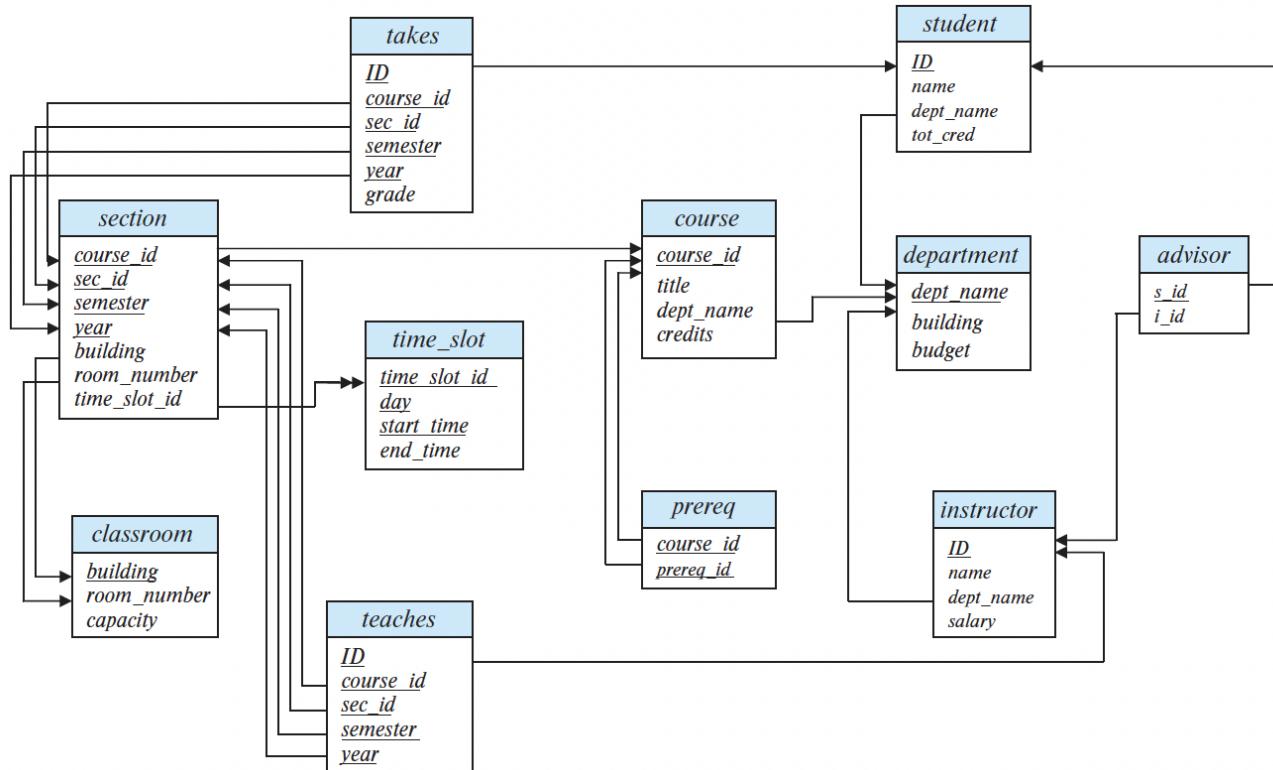


Summary--Keys

- Superkey
- Candidate key
- Primary key
- Foreign key

Schema Diagram

- A database schema, along with primary key and foreign-key constraints, can be depicted by **schema diagrams**



Agenda

- Relational Model
 - Relation, tuple, attribute
 - Schema and instance
 - Keys: Superkey, candidate key, primary key, foreign key
- Relational Algebra



Relational Algebra

- Relational Algebra: A language consisting of a set of operations that take one or two relations as input and produce a new relation as their result.
- Six basic operators
 - select: σ
 - project: Π
 - union: \cup
 - set difference: $-$
 - Cartesian product: \times
 - rename: ρ

Select Operation

- The **select** operation selects tuples that satisfy a given predicate.
- Notation: $\sigma_p(r)$
 - p is called the **selection predicate**
- Example: select those tuples of the *instructor* relation where the instructor is in the “Physics” department.
- Query

$$\sigma_{dept_name = "Physics"}(instructor)$$

Select Operation

- Example: select those tuples of the *instructor* relation where the instructor is in the “Physics” department.
- $\sigma_{dept_name = "Physics"}(instructor)$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Result

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
33456	Gold	Physics	87000

Select Operation

- Example: select the instructor named “Kim”.
- $\sigma_{name='Kim'}(instructor)$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Result

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
98345	Kim	Elec. Eng.	80000

Select Operation

- We allow comparisons using $=, \neq, >, \geq, <, \leq$ in the selection predicate.
- Example: select the instructors whose salary is greater than 80000
- $\sigma_{\text{salary} > 80000}(\text{instructor})$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Result

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
12121	Wu	Finance	90000
22222	Einstein	Physics	95000
33456	Gold	Physics	87000
83821	Brandt	Comp. Sci.	92000

Select Operation

- We can combine several predicates into a larger predicate by using the connectives: \wedge (and), \vee (or), \neg (not)
- Example: select the Physics instructors whose salary is greater than 90000
- $\sigma_{dept_name = "Physics" \wedge salary > 90000} (instructor)$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Result

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000

Project Operation

- A unary operation that returns its argument relation, with certain attributes left out.
- Notation:

$$\prod_{A_1, A_2, A_3 \dots A_k} (r)$$

where A_1, A_2, \dots, A_k are attribute names and r is a relation name.

- The result is defined as the relation of k columns obtained by erasing the columns that are not listed
- Duplicate rows removed from result, since relations are sets



Project Operation

- Example: return the *dept_name* attribute of the *instructor*.
- $\prod_{instructor} dept_name$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Result

<i>dept_name</i>
Comp. Sci.
Finance
Music
Physics
History
Biology
Elec. Eng.

Composition of Relational Operations

- Consider the query -- Find the names of all instructors in the Physics department.
- $\prod_{name}(\sigma_{dept_name = "Physics"}(instructor))$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Result

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
33456	Gold	Physics	87000



<i>name</i>
Einstein
Gold

Composition of Relational Operations

- Query -- Find the names of all instructors in the Physics department.
- Think about this: $\sigma_{dept_name = "Physics"} (\Pi_{name}(instructor))$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Wrong

- The projection goes before the selection here
- Since the projection eliminates “*dept_name*”, the selection cannot be performed

Exercise 1

- Write the following query in relational algebra and think about the result.
- Query -- Find the IDs of all instructors whose salary is less than 82000.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000



Exercise 1--Solution

- Query -- Find the IDs of all instructors whose salary is less than 82000.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

$$\Pi_{ID}(\sigma_{\text{salary} < 82000}(\text{instructor}))$$



Cartesian-Product Operation

- The Cartesian-product operation (denoted by X) allows us to combine information from any two relations.
- Example: the Cartesian product of the relations *instructor* and *teaches* is written as:

instructor X teaches

- *instructor (ID, name, dept_name, salary)*
- *teaches (ID, course_id, sec_id, semester, year)*

Cartesian-Product Operation

instructor X teaches

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

instructor

<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	CS-101	1	Fall	2017
10101	CS-315	1	Spring	2018
10101	CS-347	1	Fall	2017
12121	FIN-201	1	Spring	2018
15151	MU-199	1	Spring	2018
22222	PHY-101	1	Fall	2017
32343	HIS-351	1	Spring	2018
45565	CS-101	1	Spring	2018
45565	CS-319	1	Spring	2018
76766	BIO-101	1	Summer	2017
76766	BIO-301	1	Summer	2018
83821	CS-190	1	Spring	2017
83821	CS-190	2	Spring	2017
83821	CS-319	2	Spring	2018
98345	EE-181	1	Spring	2017

teaches



Cartesian-Product Operation

<i>instructor.ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>teaches.ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2017
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2017
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2017
...
...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2017
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2018
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2017
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2018
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2018
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2017
...
...
15151	Mozart	Music	40000	10101	CS-101	1	Fall	2017
15151	Mozart	Music	40000	10101	CS-315	1	Spring	2018
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2017
15151	Mozart	Music	40000	12121	FIN-201	1	Spring	2018
15151	Mozart	Music	40000	15151	MU-199	1	Spring	2018
15151	Mozart	Music	40000	22222	PHY-101	1	Fall	2017
...
...
22222	Einstein	Physics	95000	10101	CS-101	1	Fall	2017
22222	Einstein	Physics	95000	10101	CS-315	1	Spring	2018
22222	Einstein	Physics	95000	10101	CS-347	1	Fall	2017
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2018
22222	Einstein	Physics	95000	15151	MU-199	1	Spring	2018
22222	Einstein	Physics	95000	22222	PHY-101	1	Fall	2017
...
...



Theta Join Operation

- The Cartesian-Product
 - instructor X teaches
 - associates every tuple of instructor with every tuple of teaches.
 - Most of the resulting rows have information about instructors who did NOT teach a particular course.
- To get only those tuples of “instructor X teaches” that pertain to instructors and the courses that they taught, we write:

$$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$$

Theta Join Operation

- The table corresponding to:

$$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$$

<i>instructor.ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>teaches.ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2017
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2017
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2018
15151	Mozart	Music	40000	15151	MU-199	1	Spring	2018
22222	Einstein	Physics	95000	22222	PHY-101	1	Fall	2017
32343	El Said	History	60000	32343	HIS-351	1	Spring	2018
45565	Katz	Comp. Sci.	75000	45565	CS-101	1	Spring	2018
45565	Katz	Comp. Sci.	75000	45565	CS-319	1	Spring	2018
76766	Crick	Biology	72000	76766	BIO-101	1	Summer	2017
76766	Crick	Biology	72000	76766	BIO-301	1	Summer	2018
83821	Brandt	Comp. Sci.	92000	83821	CS-190	1	Spring	2017
83821	Brandt	Comp. Sci.	92000	83821	CS-190	2	Spring	2017
83821	Brandt	Comp. Sci.	92000	83821	CS-319	2	Spring	2018
98345	Kim	Elec. Eng.	80000	98345	EE-181	1	Spring	2017

Theta Join Operation

- The **join** operation allows us to combine a select operation and a Cartesian-Product operation into a single operation.
- Consider relations $r(R)$ and $s(S)$
- Let “theta” be a predicate on attributes in the schema R “union” S. The join operation $r \bowtie_{\theta} s$ is defined as follows:

$$r \bowtie_{\theta} s = \sigma_{\theta} (r \times s)$$

- Thus

$$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$$

- Can equivalently be written as

instructor $\bowtie_{instructor.id = teaches.id}$ *teaches*

Theta Join Operation

- Example: Find the students whose score in quiz2 is higher than quiz1
- $quiz1 \bowtie_{quiz1.name = quiz2.name \wedge quiz2.score > quiz1.score} quiz2$

Result

<i>quiz1</i>		<i>quiz2</i>	
<i>name</i>	<i>score</i>	<i>name</i>	<i>score</i>
Anna	80	Anna	90
Bob	70	Bob	60
Cathy	90	Cathy	100
David	100	David	80

<i>quiz1.name</i>	<i>quiz1.score</i>	<i>quiz2.name</i>	<i>quiz2.score</i>
Anna	80	Anna	90
Cathy	90	Cathy	100

Theta Join Operation

- Example: For those students who have made donations, find their names, department name, and amounts of their donations
- $student \bowtie_{student.name = donation.name} donation$

<i>student</i>		<i>donation</i>	
<i>name</i>	<i>dept_name</i>	<i>name</i>	<i>amount</i>
Anna	Comp. Sci.	Alice	100
Bob	Physics	Anna	150
Cathy	Music	Christian	120
David	Physics	David	180

Result

<i>student.name</i>	<i>dept_name</i>	<i>donation.name</i>	<i>amount</i>
Anna	Comp. Sci.	Anna	150
David	Physics	David	180

Any Problem

Theta Join Operation

- Use projection operation
- $\prod_{student.name, dept_name, amount} (student \bowtie_{student.name = donation.name} donation)$

<i>student</i>		<i>donation</i>	
<i>name</i>	<i>dept_name</i>	<i>name</i>	<i>amount</i>
Anna	Comp. Sci.	Alice	100
Bob	Physics	Anna	150
Cathy	Music	Christian	120
David	Physics	David	180

Result

<i>student.name</i>	<i>dept_name</i>	<i>amount</i>
Anna	Comp. Sci.	150
David	Physics	180

Natural Join Operation

- Note 1: The join is performed based on the common attributes of the two relations
- Note 2: Each common attribute appears only once in the result
- $student \bowtie donation$

<i>student</i>		<i>donation</i>	
<i>name</i>	<i>dept_name</i>	<i>name</i>	<i>amount</i>
Anna	Comp. Sci.	Alice	100
Bob	Physics	Anna	150
Cathy	Music	Christian	120
David	Physics	David	180

Result

<i>name</i>	<i>dept_name</i>	<i>amount</i>
Anna	Comp. Sci.	150
David	Physics	180

Natural Join Operation

- $(\sigma_{dept_name = "Physics"} student) \bowtie donation$

student

<i>name</i>	<i>dept_name</i>
Anna	Comp. Sci.
Bob	Physics
Cathy	Music
David	Physics

donation

<i>name</i>	<i>amount</i>
Alice	100
Anna	150
Christian	120
David	180

Result

<i>name</i>	<i>dept_name</i>	<i>amount</i>
David	Physics	180



Exercise 2

- Write the following operations in relational algebra and think about the results
 - Cartesian-product operation
 - Theta join operation
 - $student.stu_ID = course.stu_ID$
 - Natural join operation

student

<i>stu_ID</i>	<i>stu_name</i>	<i>age</i>
1	John	20
2	Sarah	19
3	Mike	22

course

<i>cou_ID</i>	<i>cou_name</i>	<i>stu_ID</i>
101	<i>Math</i>	1
102	<i>History</i>	3
103	<i>Physics</i>	2



Exercise 2--Solution

- Cartesian-product operation
- *student X course*

<i>student.stu_ID</i>	<i>stu_name</i>	<i>age</i>	<i>cou_ID</i>	<i>cou_name</i>	<i>course.stu_ID</i>
1	John	20	101	<i>Math</i>	1
1	John	20	102	<i>History</i>	3
1	John	20	103	<i>Physics</i>	2
2	Sarah	19	101	<i>Math</i>	1
2	Sarah	19	102	<i>History</i>	3
2	Sarah	19	103	<i>Physics</i>	2
3	Mike	22	101	<i>Math</i>	1
3	Mike	22	102	<i>History</i>	3
3	Mike	22	103	<i>Physics</i>	2

Exercise 2--Solution

- Theta join operation $student \bowtie_{student.stu_ID = course.stu_ID} course$
 - $student.stu_ID = course.stu_ID$

$student.stu_ID$	stu_name	age	cou_ID	cou_name	$course.stu_ID$
1	John	20	101	Math	1
2	Sarah	19	103	Physics	2
3	Mike	22	102	History	3

- Natural join operation $student \bowtie course$

stu_ID	stu_name	age	cou_ID	cou_name
1	John	20	101	Math
2	Sarah	19	103	Physics
3	Mike	22	102	History



Union Operation

- The union operation allows us to combine two relations
- Notation: $r \cup s$
- For $r \cup s$ to be valid.
 - 1. r, s must have the same **arity** (same number of attributes)
 - 2. The attribute domains must be **compatible** (example: 2nd column of r deals with the same type of values as does the 2nd column of s)

Union Operation

- Example: find the ids of all courses taught in the Fall 2017 semester, or in the Spring 2018 semester, or in both
- $\prod_{course_id} (\sigma_{semester = "Fall"} \wedge year = 2017) (section) \cup \prod_{course_id} (\sigma_{semester = "Spring"} \wedge year = 2018) (section)$

<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>	<i>building</i>	<i>room_number</i>	<i>time_slot_id</i>
BIO-101	1	Summer	2017	Painter	514	B
BIO-301	1	Summer	2018	Painter	514	A
CS-101	1	Fall	2017	Packard	101	H
CS-101	1	Spring	2018	Packard	101	F
CS-190	1	Spring	2017	Taylor	3128	E
CS-190	2	Spring	2017	Taylor	3128	A
CS-315	1	Spring	2018	Watson	120	D
CS-319	1	Spring	2018	Watson	100	B
CS-319	2	Spring	2018	Taylor	3128	C
CS-347	1	Fall	2017	Taylor	3128	A
EE-181	1	Spring	2017	Taylor	3128	C
FIN-201	1	Spring	2018	Packard	101	B
HIS-351	1	Spring	2018	Painter	514	C
MU-199	1	Spring	2018	Packard	101	D
PHY-101	1	Fall	2017	Watson	100	A

section

Result

<i>course_id</i>
CS-101
CS-315
CS-319
CS-347
FIN-201
HIS-351
MU-199
PHY-101

Union Operation

- Example: find the persons who are either students or volunteers
- $student \cup volunteer$

<i>student</i>	
<i>name</i>	<i>age</i>
Anna	19
Bob	22
Cathy	20
David	21

<i>volunteer</i>	
<i>name</i>	<i>age</i>
Alice	23
Anna	19
Christian	24
David	21

<i>Result</i>	
<i>name</i>	<i>age</i>
Alice	23
Anna	19
Bob	22
Cathy	20
Christian	24
David	21

Union Operation

- Example: find the names of the persons who are either students or volunteers
- $\prod_{name} (student \cup volunteer)$
- $\prod_{name} (student) \cup \prod_{name} (volunteer)$

<i>student</i>	
<i>name</i>	<i>age</i>
Anna	19
Bob	22
Cathy	20
David	21

<i>volunteer</i>	
<i>name</i>	<i>age</i>
Alice	23
Anna	19
Christian	24
David	21

<i>Result</i>
<i>name</i>
Alice
Anna
Bob
Cathy
Christian
David

Union Operation

- Example: find the persons who are either students or volunteers
- $student \cup volunteer$ **wrong!**
- $\prod_{name} (student) \cup volunteer$ **correct version**

<i>student</i>		<i>volunteer</i>	<i>Result</i>
<i>name</i>	<i>age</i>	<i>name</i>	<i>name</i>
Anna	19	Alice	Alice
Bob	22	Anna	Anna
Cathy	20	Christian	Bob
David	21	David	Cathy
			Christian
			David

Set-Intersection Operation

- The set-intersection operation allows us to find tuples that are in both the input relations.
- Notation: $r \cap s$
- Assume:
 - r, s have the same arity
 - attributes of r and s are compatible

Set-Intersection Operation

- Example: find the ids of all courses taught in both the Fall 2017 and the Spring 2018 semesters
- $\prod_{course_id} (\sigma_{semester = "Fall"} \wedge year = 2017)(section) \cap \prod_{course_id} (\sigma_{semester = "Spring"} \wedge year = 2018)(section)$

course_id	sec_id	semester	year	building	room_number	time_slot_id
BIO-101	1	Summer	2017	Painter	514	B
BIO-301	1	Summer	2018	Painter	514	A
CS-101	1	Fall	2017	Packard	101	H
CS-101	1	Spring	2018	Packard	101	F
CS-190	1	Spring	2017	Taylor	3128	E
CS-190	2	Spring	2017	Taylor	3128	A
CS-315	1	Spring	2018	Watson	120	D
CS-319	1	Spring	2018	Watson	100	B
CS-319	2	Spring	2018	Taylor	3128	C
CS-347	1	Fall	2017	Taylor	3128	A
EE-181	1	Spring	2017	Taylor	3128	C
FIN-201	1	Spring	2018	Packard	101	B
HIS-351	1	Spring	2018	Painter	514	C
MU-199	1	Spring	2018	Packard	101	D
PHY-101	1	Fall	2017	Watson	100	A

section

Result

course_id
CS-101

Set-Intersection Operation

- Example: find the persons who are both students and volunteers
- $student \cap volunteer$

<i>student</i>	
<i>name</i>	<i>age</i>
Anna	19
Bob	22
Cathy	20
David	21

<i>volunteer</i>	
<i>name</i>	<i>age</i>
Alice	23
Anna	19
Christian	24
David	21

Result

<i>Result</i>	
<i>name</i>	<i>age</i>
Anna	19
David	21

Set-Intersection Operation

- Example: find the names of the persons who are both students and volunteers
- $\prod_{name} (student \cap volunteer)$
- $\prod_{name} (student) \cap \prod_{name} (volunteer)$

student	
<i>name</i>	<i>age</i>
Anna	19
Bob	22
Cathy	20
David	21

volunteer	
<i>name</i>	<i>age</i>
Alice	23
Anna	19
Christian	24
David	21

Result
<i>name</i>
Anna
David

Set-Intersection Operation

- Example: find the persons who are both students and volunteers
- $student \cap volunteer$ **wrong!**
- $\prod_{name} (student) \cap volunteer$ **correct version**

<i>student</i>	
<i>name</i>	<i>age</i>
Anna	19
Bob	22
Cathy	20
David	21

<i>volunteer</i>	
<i>name</i>	
Alice	
Anna	
Christian	
David	

Result

<i>name</i>
Anna
David

Set-Difference Operation

- The set-difference operation allows us to find tuples that are in one relation but are not in another.
- Notation $r - s$
- Set differences must be taken between compatible relations.
 - r and s must have the same arity
 - attribute domains of r and s must be compatible

Set-Difference Operation

- Example: find the ids of all courses taught in the Fall 2017 semester, but not in the Spring 2018 semester
- $\prod_{course_id} (\sigma_{semester = "Fall"} \wedge year = 2017) (section) - \prod_{course_id} (\sigma_{semester = "Spring"} \wedge year = 2018) (section)$

<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>	<i>building</i>	<i>room_number</i>	<i>time_slot_id</i>
BIO-101	1	Summer	2017	Painter	514	B
BIO-301	1	Summer	2018	Painter	514	A
CS-101	1	Fall	2017	Packard	101	H
CS-101	1	Spring	2018	Packard	101	F
CS-190	1	Spring	2017	Taylor	3128	E
CS-190	2	Spring	2017	Taylor	3128	A
CS-315	1	Spring	2018	Watson	120	D
CS-319	1	Spring	2018	Watson	100	B
CS-319	2	Spring	2018	Taylor	3128	C
CS-347	1	Fall	2017	Taylor	3128	A
EE-181	1	Spring	2017	Taylor	3128	C
FIN-201	1	Spring	2018	Packard	101	B
HIS-351	1	Spring	2018	Painter	514	C
MU-199	1	Spring	2018	Packard	101	D
PHY-101	1	Fall	2017	Watson	100	A

section

Result

<i>course_id</i>
CS-347
PHY-101

Exercise 3

- Write the following operations in relational algebra and think about the results
 - Union operation
 - Set-intersection operation
 - Set-difference operation (in *Course_A*, not in *Course_B*)

<i>Course_A</i>		<i>Course_B</i>	
<i>cou_ID</i>	<i>cou_name</i>	<i>cou_ID</i>	<i>cou_name</i>
101	Math	102	History
102	History	103	Physics
103	Physics	104	Chemistry

Exercise 3--Solution

- Union
- $Course_A \cup Course_B$
- Set-intersection
- $Course_A \cap Course_B$
- Set-difference
- (in $Course_A$, not in $Course_B$)
- $Course_A - Course_B$

<i>cou_ID</i>	<i>cou_name</i>
101	Math
102	History
103	Physics
104	Chemistry

<i>cou_ID</i>	<i>cou_name</i>
102	History
103	Physics

<i>cou_ID</i>	<i>cou_name</i>
101	Math

The Assignment Operation

- It is convenient at times to write a relational-algebra expression by assigning parts of it to temporary relation variables.
- The assignment operation is denoted by \leftarrow and works like assignment in a programming language.
- Example: Find all instructors in the Physics or Music department.

$$\text{Physics} \leftarrow \sigma_{\text{dept_name} = \text{"Physics"}}(\text{instructor})$$
$$\text{Music} \leftarrow \sigma_{\text{dept_name} = \text{"Music"}}(\text{instructor})$$
$$\text{Physics} \cup \text{Music}$$

- It can also be used to attributes
 - $\prod_{A \leftarrow B}(r)$

The Assignment Operation

- Example: find all instructors in the Physics or Music departments

$Physics \leftarrow \sigma_{dept_name = "Physics"}(instructor)$

$Music \leftarrow \sigma_{dept_name = "Music"}(instructor)$

$Physics \cup Music$

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

instructor

<i>Physics</i>			
ID	name	dept_name	salary
22222	Einstein	Physics	95000
33456	Gold	Physics	87000

<i>Music</i>			
ID	name	dept_name	salary
15151	Mozart	Music	40000

<i>Result</i>			
ID	name	dept_name	salary
22222	Einstein	Physics	95000
33456	Gold	Physics	87000
15151	Mozart	Music	40000

The Rename Operation

- The results of relational-algebra expressions do not have a name that we can use to refer to them. The rename operator, ρ , is provided for that purpose
- The expression:

$$\rho_x(E)$$

returns the result of expression E under the name x

- Another form of the rename operation:

$$\rho_{x(A1, A2, \dots An)}(E)$$

The Rename Operation

- Example: find all instructor in the Physics or Music department

 $\rho_{Physics} (\sigma_{dept_name = "Physics"} (instructor))$
 $Physics \leftarrow \sigma_{dept_name = "Physics"} (instructor)$
 $\rho_{Music} (\sigma_{dept_name = "Music"} (instructor))$
 $Music \leftarrow \sigma_{dept_name = "Music"} (instructor)$

Physics \cup *Music*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

instructor

Physics

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
33456	Gold	Physics	87000

Music

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
15151	Mozart	Music	40000

Result

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
33456	Gold	Physics	87000
15151	Mozart	Music	40000

The Rename Operation

 $\rho_{Physics}(p.ID, p.name, p.dept, p.salary) (\sigma_{dept_name = "Physics"}(instructor))$
 $\rho_{Music}(m.ID, m.name, m.dept, m.salary) (\sigma_{dept_name = "Music"}(instructor))$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Physics

<i>p.ID</i>	<i>p.name</i>	<i>p.dept</i>	<i>p.salary</i>
22222	Einstein	Physics	95000
33456	Gold	Physics	87000

Music

<i>m.ID</i>	<i>m.name</i>	<i>m.dept</i>	<i>m.salary</i>
15151	Mozart	Music	40000

instructor

Left Join Operation

- Example: For each student, find the amount of his/her donation
- *Student ⚠ donation*

<i>student</i>		<i>donation</i>	
<i>name</i>	<i>dept_name</i>	<i>name</i>	<i>amount</i>
Anna	Comp. Sci.	Alice	100
Bob	Physics	Anna	150
Cathy	Music	Christian	120
David	Physics	David	180

Result

<i>name</i>	<i>dept_name</i>	<i>amount</i>
Anna	Comp. Sci.	150
Bob	Physics	null
Cathy	Music	null
David	Physics	180

Right Join Operation

- Example: For each donator, find the department he/she is in
- *Student ⚠ donation*

<i>student</i>		<i>donation</i>	
<i>name</i>	<i>dept_name</i>	<i>name</i>	<i>amount</i>
Anna	Comp. Sci.	Alice	100
Bob	Physics	Anna	150
Cathy	Music	Christian	120
David	Physics	David	180

Result

<i>name</i>	<i>dept_name</i>	<i>amount</i>
Alice	null	100
Anna	Comp. Sci.	150
Christian	null	120
David	Physics	180

Outer Join Operation

- *Student \bowtie donation*

Result

<i>student</i>		<i>donation</i>	
<i>name</i>	<i>dept_name</i>	<i>name</i>	<i>amount</i>
Anna	Comp. Sci.	Alice	100
Bob	Physics	Anna	150
Cathy	Music	Christian	120
David	Physics	David	180

<i>name</i>	<i>dept_name</i>	<i>amount</i>
Alice	null	100
Anna	Comp. Sci.	150
Bob	Physics	null
Cathy	Music	null
Christian	null	120
David	Physics	180



Grouping and Aggregation

- Tuples with the same attribute values (for a specified list of attributes) are grouped.
- An aggregate function is applied to each group (computing one value for each group).
- Typical aggregate functions
 - count: number of tuples in a group
 - sum: sum of attribute values in a group
 - min, max, avg: minimum, maximum, average of attribute values in a group
- Notation: $\gamma_{L; F}(r)$
 - L list of attributes for grouping
 - F aggregate function

Grouping and Aggregation

- Example: determine the number of students per semester
- $\gamma_{semester; count(*)} (student)$

student

<i>name</i>	<i>semester</i>	<i>age</i>
Alice	7	24
Anna	3	20
Bob	3	19
Cathy	5	22
Christian	5	20
David	1	18

Result

<i>semester</i>	<i>count(*)</i>
7	1
3	2
5	2
1	1

Grouping and Aggregation

- Example: determine the number of students per semester and the minimum age in each group
- $\gamma_{semester; count(*), min(age)} (student)$

student

<i>name</i>	<i>semester</i>	<i>age</i>
Alice	7	24
Anna	3	20
Bob	3	19
Cathy	5	22
Christian	5	20
David	1	18

Result

<i>semester</i>	<i>count(*)</i>	<i>min(age)</i>
7	1	24
3	2	19
5	2	20
1	1	18

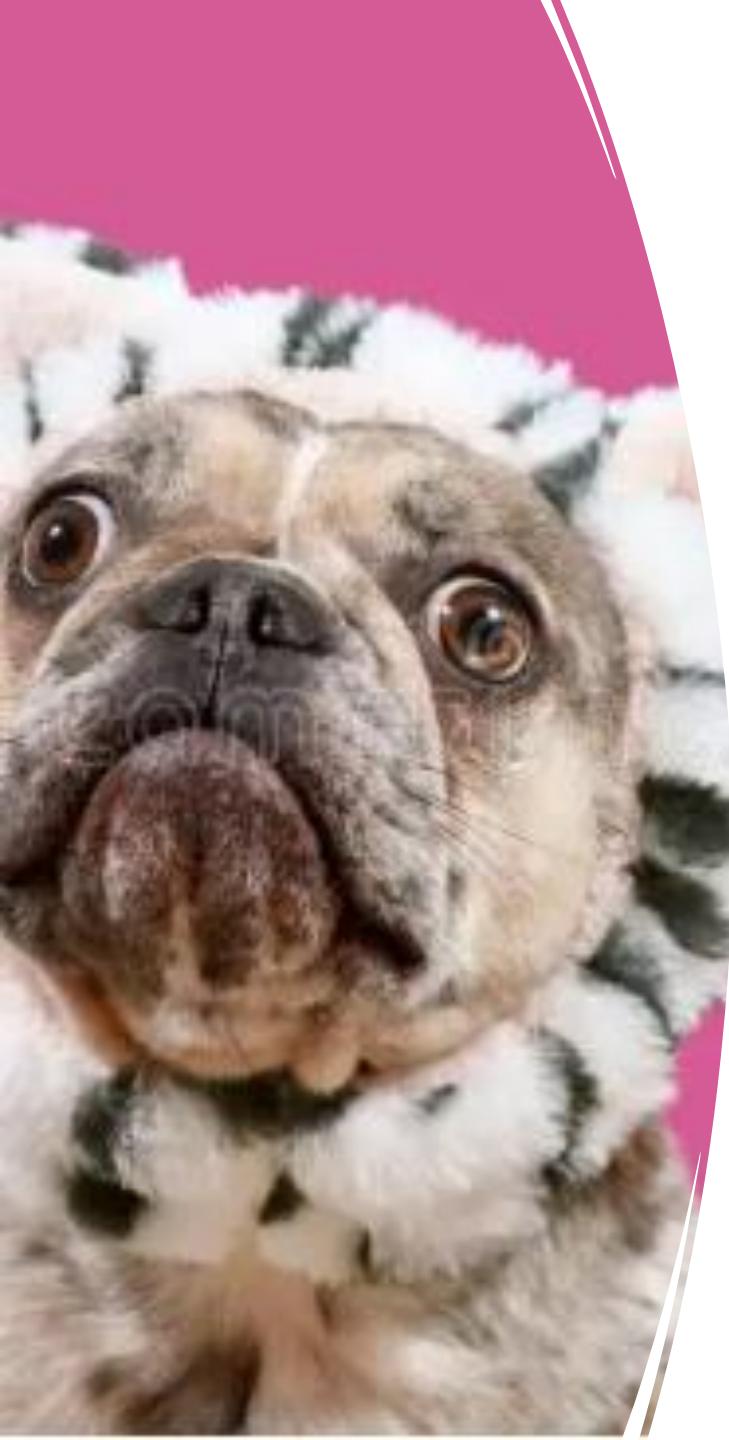
Equivalent Queries

- There is more than one way to write a query in relational algebra.
- Example: find instructors in the Physics department with salary greater than 90,000
- Query 1
 - $\sigma_{dept_name = "Physics"} \wedge salary > 90,000 (instructor)$
- Query 2
 - $\sigma_{dept_name = "Physics"} (\sigma_{salary > 90,000} (instructor))$
- The two queries are not identical; they are, however, equivalent -- they give the same result on any database.



Summary

- Relational Model
 - Relation, tuple, attribute
 - Schema and instance
 - Keys: Superkey, candidate key, primary key, foreign key
- The Relational Algebra
 - select: σ
 - project: Π
 - Cartesian product: \times
 - join: \bowtie
 - union: \cup
 - set-intersection: \cap
 - set-difference: $-$
 - rename: ρ



Next Lecture

- SQL
 - Explain and use the SQL data model
 - Create non-trivial database tables
 - Modify non-trivial database tables
 - Create non-trivial SQL statements