



AALBORG  
UNIVERSITY

# Database System (SW5)

## 10. Transaction and Concurrency Control

**Tiantian Liu**  
Department of Computer Science  
Aalborg University  
Fall 2025



# Motivation

- Transaction boundaries are an important part of system design
- Users think in transactions
- Simple to handle multiple users



# Learning Goals

- Understanding the transaction concept
- Understanding serializability
- Understand and use lock-based concurrency control
- Understand and use two-phase locking



# Agenda

- Transaction
  - Transaction concept
  - Schedules
  - Conflict serializability
- Concurrency Control
  - Lock-based synchronization
  - Two-phase locking (2PL)
  - Deadlock



# Transaction Concept

- Transaction to transfer 50 from account A to account B:
  - 1. read(A)
  - 2. A := A – 50
  - 3. write(A)
  - 4. read(B)
  - 5. B := B + 50
  - 6. write(B)
- Two main issues to deal with:
  - Failures of various kinds, such as hardware failures and system crashes
  - Concurrent execution of multiple transactions



# ACID Properties

- A **transaction** is a unit of program execution that accesses and possibly updates various data items.
- To preserve the integrity of data the database system must ensure:

**Atomicity.** Either all operations of the transaction are properly reflected in the database or none are.

**Consistency.** Execution of a transaction in isolation preserves the consistency of the database.

**Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions.

**Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



# Operations on Transactions

- Begin of Transaction
  - Represents the beginning of a transaction, i.e., all following statements together form a transaction.
  - In SQL      **BEGIN**;
- Commit
  - Represents the end of a transaction, i.e., all changes are made persistent and visible to others.
  - In SQL      **COMMIT**;
- Rollback or Abort
  - Causes a transaction to roll back, i.e., all changes are undone/discard.
  - In SQL      **ROLLBACK**;



# Savepoints

Long running transactions can specify savepoints.

- **SAVEPOINT** savepoint name;
  - Defines a point/state within a transaction
  - A transaction can be **rolled back partially** back up to the savepoint.
- 
- **ROLLBACK TO** <savepoint name>:
  - rolls the active transaction back to the savepoint <savepoint name>



# Example

```
BEGIN;  
INSERT INTO tab VALUES...  
SAVEPOINT A;  
INSERT INTO tab VALUES...  
SAVEPOINT B;  
SELECT * FROM tab;  
ROLLBACK TO A;  
SELECT * FROM tab;  
...
```



# Termination of Transactions

There are three possibilities for a transaction to terminate:

- Successful termination by **commit**
- Unsuccessful termination by **abort**
- Unsuccessful termination by a **failure**



# How do DBMSs support transactions

The two most important components of transaction management are

- Multi-user synchronization (isolation)
  - Concurrency allows for high throughput
  - Serializability
- Recovery (atomicity and durability)
  - Roll back partially executed transactions
  - Re-executing transactions after failures
  - Guaranteeing persistence of transactional updates



# Concurrent Executions


- Multiple transactions are allowed to run concurrently in the system.
- Advantages are:
  - Increased processor and disk utilization, leading to better transaction throughput
    - E.g., one transaction can be using the CPU while another is reading from or writing to the disk
  - Reduced average response time for transactions: short transactions need not wait behind long ones.

# Concurrent Executions

- Affects the "I" in ACID.
- The execution of multiple transactions  $T_1$ ,  $T_2$ , and  $T_3$
- (a) in a single-user environment



- (b) in a (concurrent) multi-user environment



# Potential Problems

- Lost updates (overwriting updates)

Step	T <sub>1</sub>	T <sub>2</sub>
1	read(A, a <sub>1</sub> )	
2	a <sub>1</sub> := a <sub>1</sub> - 300	
3		read(A, a <sub>2</sub> )
4		a <sub>2</sub> := a <sub>2</sub> * 1.03
5		<b>write(A, a<sub>2</sub>)</b>
6	<b>write(A, a<sub>1</sub>)</b>	
7	read(B, b <sub>1</sub> )	
8	b <sub>1</sub> := b <sub>1</sub> + 300	
9		write(B, b <sub>1</sub> )

# Potential Problems

- Dirty read (dependency on non-committed updates)

Step	T <sub>1</sub>	T <sub>2</sub>
1	read(A, a <sub>1</sub> )	
2	a <sub>1</sub> := a <sub>1</sub> - 300	
3	<b>write(A, a<sub>1</sub>)</b>	
4		read(A, a <sub>2</sub> )
5		a <sub>2</sub> := a <sub>2</sub> * 1.03
6		<b>write(A, a<sub>2</sub>)</b>
7	read(B, b <sub>1</sub> )	
8	...	
9	abort	



# Potential Problems

- Non-repeatable read (dependency on other updates)

T <sub>1</sub>	T <sub>2</sub>
update account set balance = 42000 where accountID = 123	select sum (balance) from account

# Potential Problems

- Phantom problem (dependency on new/deleted tuples)

$T_1$	$T_2$
Insert into account values (C, 1000, ...)	select sum (balance) from account  select sum (balance) from account



# Agenda

- Transaction
  - Transaction concept
  - Schedules
  - Conflict serializability
- Concurrency Control
  - Lock-based synchronization
  - Two-phase locking (2PL)
  - Deadlock



# Schedules

- Schedule – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
- A transaction that successfully completes its execution will have a **commit** instruction as the last statement
- A transaction that fails to successfully complete its execution will have an **abort** instruction as the last statement



# Schedules

- Serial schedule
  - The operations of the transactions are executed sequentially with no overlap in time.
- Concurrent schedule
  - The operations of the transactions are executed with overlap in time.
- Valid schedule
  - A schedule is valid if the result of its execution is "correct".

# Schedule 1

- Let  $T_1$  transfer 50 from A to B, and  $T_2$  transfer 10% of the balance from A to B.
- A serial schedule in which  $T_1$  is followed by  $T_2$ :

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

# Schedule 2

- A serial schedule where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
<pre>read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) commit</pre>	<pre>read (A) temp := A * 0.1 A := A - temp write (A) read (B) B := B + temp write (B) commit</pre>

# Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is not a serial schedule, but it is equivalent to Schedule 1

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ )	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ )
read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

- In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.

# Schedule 4

- The following concurrent schedule does not preserve the value of  $(A + B)$ .

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ )
write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	$B := B + temp$ write ( $B$ ) commit



# Notion of Correctness

- Definition D1: A concurrent execution of transactions must leave the database in a consistent state.
- Definition D2: Concurrent execution of transactions must be (result) equivalent to some serial execution of the transactions.



# Agenda

- Transaction
  - Transaction concept
  - Schedules
  - Conflict serializability
- Concurrency Control
  - Lock-based synchronization
  - Two-phase locking (2PL)
  - Deadlock



# Serializability

- Basic Assumption – Each transaction preserves database consistency.
  - Thus, serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule.
- Different forms of schedule equivalence give rise to the notions of:
  - 1. Conflict serializability
  - 2. View serializability



# Simplified View of Transactions

- We ignore operations other than **read** and **write** instructions
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.

# Conflicting Instructions

- Instructions  $I_1$  and  $I_2$  of transactions  $T_1$  and  $T_2$  respectively, conflict if and only if there exists some item  $Q$  accessed by both  $I_1$  and  $I_2$ , and at least one of these instructions wrote  $Q$ .
  1.  $I_1 = \text{read}(Q), I_2 = \text{read}(Q)$ .  $I_1$  and  $I_2$  don't conflict.
  2.  $I_1 = \text{read}(Q), I_2 = \text{write}(Q)$ . They conflict.
  3.  $I_1 = \text{write}(Q), I_2 = \text{read}(Q)$ . They conflict
  4.  $I_1 = \text{write}(Q), I_2 = \text{write}(Q)$ . They conflict
- If  $I_1$  and  $I_2$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

# Conflicts--Example

- Conflicts between pairs of transactions ( $T_1$  and  $T_2$ ) and their instructions.

$T_1$	$T_2$
read(X, x)	
	write(X, x)

**Conflict**

$T_1$	$T_2$
write(X, x)	
	read(X, x)

**Conflict**

$T_1$	$T_2$
write(X, x)	
	write(X, x)

**Conflict**

$T_1$	$T_2$
read(X, x)	
	read(X, x)

**No conflict**

# Conflicts--Example

- Conflicts between pairs of transactions ( $T_1$  and  $T_2$ ) and their instructions.

$T_1$	$T_2$
read(X, x)	write(Y, y)

No conflict

$T_1$	$T_2$
write(X, x)	read(Y, y)

No conflict

$T_1$	$T_2$
write(X, x)	write(Y, y)

No conflict

$T_1$	$T_2$
read(X, x)	read(Y, y)

No conflict

# Conflict Serializability

- If a schedule  $S_1$  can be transformed into a schedule  $S_2$  by a series of swaps of non-conflicting instructions, we say that  $S_1$  and  $S_2$  are **conflict equivalent**.
- We say that a schedule  $S$  is **conflict serializable** if it is **conflict equivalent** to a serial schedule

# Conflict Serializability (Cont.)

- Let  $I$  and  $J$  be consecutive instructions of a schedule  $S_1$  of multiple transactions.
- If  $I$  and  $J$  do not conflict, we can swap their order to produce a new schedule  $S_2$ .
- The instructions appear in the same order in  $S_1$  and  $S_2$ , except for  $I$  and  $J$ , whose order does not matter.
- $S_1$  and  $S_2$  are termed conflict equivalent schedules.

# Conflict Serializability--Example

$T_1$	$T_2$
read ( $A$ ) write ( $A$ )	read ( $A$ ) write ( $A$ )
read ( $B$ ) write ( $B$ )	read ( $B$ ) write ( $B$ )

Schedule 3

$T_1$	$T_2$
read ( $A$ ) write ( $A$ ) read ( $B$ ) write ( $B$ )	read ( $A$ ) write ( $A$ ) read ( $B$ ) write ( $B$ )
	read ( $A$ ) write ( $A$ ) read ( $B$ ) write ( $B$ )

Schedule 6

# Conflict Serializability-Example

- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read ( $Q$ )	
write ( $Q$ )	write ( $Q$ )

- We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .

# Conflict Serializability

- As the transformation shows, the initial concurrent schedule is conflict equivalent to a serial schedule and is therefore conflict serializable.

T<sub>1</sub>: R(A), W(A), R(B), W(B), Commit

T<sub>2</sub>: R(C), W(C), R(A), W(A), Commit



T<sub>1</sub>: R(A), W(A), R(B), W(B), Commit

T<sub>2</sub>: R(C), W(C), R(A), W(A), Commit



T<sub>1</sub>: R(A), W(A), R(B), W(B), Commit

T<sub>2</sub>: R(C), W(C), R(A), W(A), Commit

# Exercise 1

- The following schedules are conflict serializable or not:

schedule $S_A$	
$T_1$	$T_2$
read(Y, y)	read(X, x) write(X, x)
write(Y, y)	

schedule $S_B$	
$T_3$	$T_4$
read(X, x)	read(X, x) write(X, x)
write(X, x)	

schedule $S_C$	
$T_5$	$T_6$
read(X, x)	read(X, x)
write(X, x)	

schedule $S_D$	
$T_7$	$T_8$
read(X, x)	write(X, x)
write(X, x)	

# Exercise 1--Solution

- The following schedules are conflict serializable or not:

schedule $S_A$	
$T_1$	$T_2$
read(Y, y)	read(X, x) write(X, x)
write(Y, y)	

conflict serializable

schedule $S_B$	
$T_3$	$T_4$
read(X, x)	read(X, x) write(X, x)
write(X, x)	

Not conflict serializable

schedule $S_C$	
$T_5$	$T_6$
read(X, x)	read(X, x)
write(X, x)	


conflict serializable

schedule $S_D$	
$T_7$	$T_8$
read(X, x)	write(X, x)
write(X, x)	

Not conflict serializable




# Conflict/Precedence Graph

- We construct a directed graph (conflict/precedence graph) for a schedule involving a set of transactions.
- Given a schedule for a set of transactions  $T_1, T_2, \dots, T_n$ 
  - The vertices of the conflict graph are the transaction identifiers.
  - An edge from  $T_i$  to  $T_j$  denotes that the two transactions are conflicting, with  $T_i$  making the relevant access earlier.
  - Sometimes the edge is labelled with the item involved in the conflict.




# Test for Conflict Serializability

- A schedule is conflict serializable if and only if its precedence graph is **acyclic**.
- Intuitively, a conflict between two transactions forces an execution order between them (topological sorting)




# Conflict Graph--Example

$T_{10}$	$T_{11}$	$T_{12}$	$T_{13}$	$T_{14}$
read(Y, y) read(Z, z)	read(X, x)			
	read(Y, y) write(Y, y)	read(Z, z) write(Z, z)		read(V, v) read(W, w) write(W, w)
read(T, t)			read(Y, y) write(Y, y) read(Z, z) write(Z, z)	
read(U, u)				



# Conflict Graph--Example

$T_{10}$	$T_{11}$	$T_{12}$	$T_{13}$	$T_{14}$
read(Y, y) read(Z, z)	read(X, x)			
	read(Y, y) write(Y, y)			read(V, v) read(W, w) write(W, w)
read(T, t)		read(Z, z) write(Z, z)		
read(U, u)			read(Y, y) write(Y, y) read(Z, z) write(Z, z)	




- Which of the following are conflict equivalent serial schedules?
- $T_{10}, T_{11}, T_{12}, T_{13}$ , and  $T_{14}$  Yes
- $T_{14}, T_{10}, T_{12}, T_{11}$ , and  $T_{13}$  Yes
- $T_{14}, T_{13}, T_{12}, T_{11}$ , and  $T_{10}$  No

# Example

- Draw the precedence graphs of the following schedule:
  - Decide if the schedule is conflict serializable. If it is conflict serializable, give one example of a conflict equivalent serial schedule?

$T_1:$	$R(A), W(A),$	$R(B), W(B)$
$T_2:$	$W(A),$	$R(B), W(B)$
$T_3:$	$R(A),$	$W(A)$



Conflict equivalent serial schedule:  $T_1, T_2, T_3$

# Exercise 2

- Draw the precedence graphs of the following schedules:
  - Decide if the schedule is conflict serializable. If it is conflict serializable, give one example of a conflict equivalent serial schedule?

T<sub>1</sub>: R(A), W(A),                                    R(B), W(B)

T<sub>2</sub>:    R(C), W(A)

T<sub>3</sub>:    R(B),    W(C)

T<sub>4</sub>: R(A),    W(A)

T<sub>5</sub>:    W(A)

T<sub>6</sub>:    W(A)

# Exercise 2--Solution

- Draw the precedence graphs of the following schedules:
  - Decide if the schedule is conflict serializable. If it is conflict serializable, give one example of a conflict equivalent serial schedule?


T<sub>1</sub>: R(A), W(A),

R(B), W(B)

T<sub>2</sub>: R(C), W(A)

T<sub>3</sub>: R(B),

W(C)



Not Conflict Serializable


# Exercise 2--Solution

- Draw the precedence graphs of the following schedules:
  - Decide if the schedule is conflict serializable. If it is conflict serializable, give one example of a conflict equivalent serial schedule?

$T_4: R(A), \quad W(A)$

$T_5: \quad W(A)$

$T_6: \quad W(A)$



Not Conflict Serializable



# Agenda

- Transaction
  - Transaction concept
  - Schedules
  - Conflict serializability
- Concurrency Control
  - Lock-based synchronization
  - Two-phase locking (2PL)
  - Deadlock

# Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item.
- Data items can be locked in two modes :
  - 1. **Exclusive (X)** mode. X-lock is requested using **lock-X** instruction.
  - 2. **Shared (S)** mode. S-lock is requested using **lock-S** instruction.
- Operations on locks
  - **lock-S(Q)**: set shared lock on data item Q. Allow **multiple** transactions to **read** Q simultaneously.
  - **lock-X(Q)**: set exclusive lock on data item Q. Ensure only **one** transaction can **read** or **write** Q at a time.
  - **unlock(Q)**: release lock on data item Q.

# Lock-Based Protocols (Cont.)

- Lock-compatibility matrix

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item
- But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item

# Example

- Example of a transaction performing locking:

$T_2$ : **lock-S(A);**

**read (A);**

**unlock(A);**

**lock-S(B);**

**read (B);**

**unlock(B);**

**display(A+B)**

- Locking as above is **not sufficient** to guarantee serializability

# Example

- $T_1$  transfers 50 kr. from account B to account A.
- $T_2$  displays the total amount of money in accounts A and B.
- Initially A = 100 and B = 200

Result of the serial execution:  
 $T_2$  will display 300  
Using this schedule:  
 $T_2$  will display 250

$T_1$	$T_2$
lock-X(B)	
read(B, b)	
$b \leftarrow b - 50$	lock-S(A)
write(B, b)	read(A, a)
unlock(B)	unlock(A)
lock-X(A)	lock-S(B)
read(A, a)	read(B, b)
$a \leftarrow a + 50$	unlock(B)
write(A, a)	display(a + b)
display(a + b)	
unlock(A)	

# Problems with early unlocking

- $T_1$  transfers 50 kr. from account B to account A.
- $T_2$  displays the total amount of money in accounts A and B.
- Initially A = 100 and B = 200

Early unlocking can cause incorrect results (non-Serializable schedules) but allows for a higher degree of concurrency.

$T_1$	$T_2$
lock-X(B)	
read(B, b)	
$b \leftarrow b - 50$	lock-S(A)
write(B, b)	read(A, a)
unlock(B)	unlock(A)
lock-X(A)	lock-S(B)
read(A, a)	read(B, b)
$a \leftarrow a + 50$	unlock(B)
write(A, a)	display(a + b)
unlock(A)	

# Problems with late unlocking

- $T_1$  transfers 50 kr. from account B to account A.
- $T_2$  displays the total amount of money in accounts A and B.
- Initially A = 100 and B = 200

Late unlocking avoids non-serializable schedules.

But it increases the chances of deadlocks.

$T_1$	$T_2$
lock-X(B)	
read(B, b)	
$b \leftarrow b - 50$	lock-S(A)
write(B, b)	read(A, a)
unlock(B)	unlock(A)
lock-X(A)	lock-S(B)
read(A, a)	read(B, b)
$a \leftarrow a + 50$	unlock(B)
write(A, a)	display(a + b)
unlock(B)	unlock(A)
unlock(A)	unlock(B)




# Agenda

- Transaction
  - Transaction concept
  - Schedules
  - Conflict serializability
- Concurrency Control
  - Lock-based synchronization
  - Two-phase locking (2PL)
  - Deadlock

# The Two-Phase Locking Protocol

- A protocol which ensures conflict-serializable schedules.
- Phase 1: Growing Phase
  - Transaction may obtain locks
  - Transaction may not release locks
- Phase 2: Shrinking Phase
  - Transaction may release locks
  - Transaction may not obtain locks



# 2PL: yes or no?

---

 $T_i$ 

---

lock-X(A)

lock-X(B)

lock-X(C)

unlock(A)

unlock(B)

unlock(C)

---

 $T_i$ 

---

lock-X(A)

lock-X(B)

unlock(B)

lock-X(C)

unlock(A)

unlock(C)

yes

no



# Agenda

- Transaction
  - Transaction concept
  - Schedules
  - Conflict serializability
- Concurrency Control
  - Lock-based synchronization
  - Two-phase locking (2PL)
  - Deadlock

# Deadlock Handling

- System is **deadlocked** if there is a set of transactions such that every transaction in the set is waiting for another transaction in the set.

$T_3$	$T_4$
lock-X( $B$ ) read( $B$ ) $B := B - 50$ write( $B$ )  lock-X( $A$ )	lock-S( $A$ ) read( $A$ ) lock-S( $B$ )




# Deadlock Handling

- Deadlock prevention protocols ensure that the system will never enter into a deadlock state. Some prevention strategies:
  - Pre-declaration
  - Graph-based protocol


# Deadlock Detection

- Wait-for graph
  - Vertices: transactions
  - Edge from  $T_i \rightarrow T_j$ : if  $T_i$  is waiting for a lock held in conflicting mode by  $T_j$
- The system is in a deadlock state if and only if the wait-for graph has a cycle.
- Invoke a deadlock-detection algorithm periodically to look for cycles.

# Wait-for Graph




Wait-for graph without a cycle



Wait-for graph with a cycle

# Wait-for Graph Example

$T_1$	$T_2$	$T_3$	$T_4$
	lock-X(A)		
lock-X(B)			
lock-X(C)			
lock-X(D)		lock-X(B)	
	lock-X(D)		lock-X(E)
lock-X(E)			lock-X(A)



- Rollback of one or multiple involved transactions to release the deadlock

# Exercise 3


- Draw the wait-for graphs of the following schedules:
  - Decide if there is a deadlock

$T_1$	$T_2$	$T_3$
	lock-X(A)	
lock-X(A)		lock-X(B) lock-X(C)
lock-X(B)		lock-X(D)

# Exercise 3--Solution

- Draw the wait-for graphs of the following schedules:
  - Decide if there is a deadlock

$T_1$	$T_2$	$T_3$
	lock-X(A)	
lock-X(A)		lock-X(B) lock-X(C)
lock-X(B)		lock-X(D)



# Deadlock Recovery

- When deadlock is detected :
  - Some transaction will have to rolled back (made a victim) to break deadlock cycle.
  - Rollback -- determine how far to roll back transaction
    - Total rollback: Abort the transaction and then restart it.
    - Partial rollback: Roll back victim transaction only as far as necessary to release locks that another transaction in cycle is waiting for



# Summary

- Transaction
  - Transaction concept
  - Schedules
  - Conflict serializability
- Concurrency Control
  - Lock-based synchronization
  - Two-phase locking (2PL)
  - Deadlock



# Next Lecture

- Lecture Session
  - Recovery
    - Understanding basic logging algorithms
    - Understanding the importance of atomicity and durability
  - Summary
- Exercise Session
  - Mock exam
  - Q & A