



AALBORG
UNIVERSITY

Database System (SW5)

11. Recovery

Tiantian Liu
Department of Computer Science
Aalborg University
Fall 2025



AALBORG
UNIVERSITY

Motivation and Learning Goals

- Motivation
 - We want to preserve consistency and availability even in the case of failures.
- Learning Goals
 - Understanding basic logging algorithms
 - Understanding the importance of atomicity and durability



Recovery

- "Problems" with transactions
 - Atomicity: Transactions may abort (rollback)
 - Durability: What if a DBMS crashes?
- The DBMS ensures that a transaction
 - Either completes and has a permanent result (committed transaction)
 - Or has no effect at all on the database (aborted transaction).
- The role of the recovery component is to ensure **atomicity** and **durability** of transactions in the presence of system failures.



Failure Classification

- Transaction failure (failure of a not yet committed transaction)
 - Undo the changes of the transaction
- System crash (failure with main memory loss)
 - Changes of committed transactions must be preserved
 - Changes of all non-committed transactions need to be undone
- Disk failure (failure with hard disk loss)
 - Recovery based on archives/dumps



Data Access

- **Physical blocks** are those blocks residing on the disk.
- **Buffer blocks** are the blocks residing temporarily in main memory.
- Block movements between disk and main memory are initiated through the following two operations:
 - Input (B) transfers the physical block B to main memory.
 - Output (B) transfers the buffer block B to the disk, and replaces the appropriate physical block there.
- We assume, for simplicity, that each data item fits in, and is stored inside, a single block.

The WAL rule for log-based recovery

- WAL (Write Ahead Logging)
 - Before a transaction enters the commit state, "all its" log entries have to be written to stable storage, incl. the commit log entry
 - Before a modified page (or block) in main memory can be written to the database (non-volatile storage), "all its" log entries have to be written to stable storage

Structure of a log entry (log record)


- [TID, DID, old, new]
 - TID identifier of the transaction that caused the update
 - DID data item identifier: location on disk (page, block)
 - old value of the data item before the update
 - new value of the data item after the update
- Additional entries
 - start Transaction TID has started [TID start]
 - commit Transaction TID has committed [TID commit]
 - abort Transaction TID has aborted [TID abort]

Log entry example


T_1	T_2	
begin		[TID, DID, old, new]
read(B, b)		
$b \leftarrow b + 100$		
write(B, b)		[T1 start]
commit		[T1, B, 300, 400]
	begin	
	read(D, d)	[T1 commit]
	$d \leftarrow d - 10$	
	write(D, d)	[T2 start]
	read(E, e)	[T2, D, 530, 520]
	$e \leftarrow e - 20$	
	write(E, e)	[T2, E, 70, 50]
	commit	[T2 commit]

Log-based recovery

- Operations to recover from failures
 - Redo: perform the changes to the database again
 - Undo: restore database to state prior to execution
- To recover from a failure
 - Redo results for committed transactions
 - Undo changes of transactions that did not commit



Redo T_1 and T_2
Undo T_3



Redo T_4
Undo T_5



The phases of recovery

- Redo (repeat history)
 - Forward scan through the log
 - Repeat all updates in the same order as in the log file
 - Determine "undo" transactions
 - [TID start] add TID to the "undo list"
 - [TID abort] or [TID commit] remove TID from the "undo list"
- Undo (rollback) all transactions in the "undo list"
 - Backward scan through the log
 - Undo all updates of transactions in the "undo list": create a compensating log record
 - For a [TID start] record of a transaction TID in the "undo list", add a [TID abort] record to the log file, remove TID from the "undo list"
 - Stop, when "undo list" is empty



Compensation log records

- [TID, DID, old]
 - Created to undo (compensate) the changes of [TID, DID, old, new]
 - Redo-only log record
 - Can also be used to rollback a transaction during normal operation

Example

Database

A	100
B	200
C	50
D	60
E	120

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	200
C	50
D	60
E	120

undo list
{}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]



Example

Phase 1 (redo)

Database

A	100
B	200
C	50
D	60
E	120

undo list
{T1}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	200
C	50
D	60
E	120

undo list
{T1}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	400
C	50
D	60
E	120

undo list
{T1}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	400
C	50
D	60
E	120

undo list
{T1}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	400
C	100
D	60
E	120

undo list
{T1}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	400
C	100
D	60
E	120

undo list
{T1}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	400
C	100
D	60
E	120

undo list
{T1, T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	400
C	100
D	60
E	120

undo list
 $\{T1, T2\}$

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	400
C	100
D	60
E	480

undo list
 $\{T1, T2\}$

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	100
B	400
C	100
D	60
E	480

undo list
 $\{T1, T2\}$

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	320
B	400
C	100
D	60
E	480

undo list
 $\{T1, T2\}$

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	320
B	400
C	100
D	60
E	480

undo list
 $\{T1, T2\}$

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	320
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	320
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]



Example

Phase 1 (redo)

Database

A	520
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	520
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 1 (redo)

Database

A	520
B	400
C	100
D	530
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 2 (undo)

Database

A	520
B	400
C	100
D	530
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 2 (undo)

Database

A	520
B	400
C	100
D	530
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

Example

Phase 2 (undo)

Database

A	520
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]
[T1, B, 200, 400]
[T1, C, 50, 100]
[T2 start]
[T2, E, 120, 480]
[T1, A, 100, 320]
[T1 commit]
[T2, A, 320, 520]
[T2, D, 60, 530]
[T2, D, 60]

Example

Phase 2 (undo)

Database

A	520
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]
[T1, B, 200, 400]
[T1, C, 50, 100]
[T2 start]
[T2, E, 120, 480]
[T1, A, 100, 320]
[T1 commit]
[T2, A, 320, 520]
[T2, D, 60, 530]
[T2, D, 60]

Example

Phase 2 (undo)

Database

A	320
B	400
C	100
D	60
E	480

undo list

{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

[T2, D, 60]

[T2, A, 320]

Example

Phase 2 (undo)

Database

A	320
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

[T2, D, 60]

[T2, A, 320]

Example

Phase 2 (undo)

Database

A	320
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

[T2, D, 60]

[T2, A, 320]

Example

Phase 2 (undo)

Database

A	320
B	400
C	100
D	60
E	480

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

[T2, D, 60]

[T2, A, 320]

Example

Phase 2 (undo)

Database

A	320
B	400
C	100
D	60
E	120

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

[T2, D, 60]

[T2, A, 320]

[T2, E, 120]

Example

Phase 2 (undo)

Database

A	320
B	400
C	100
D	60
E	120

undo list
{T2}

Log records

[T1 start]

[T1, B, 200, 400]

[T1, C, 50, 100]

[T2 start]

[T2, E, 120, 480]

[T1, A, 100, 320]

[T1 commit]

[T2, A, 320, 520]

[T2, D, 60, 530]

[T2, D, 60]

[T2, A, 320]

[T2, E, 120]

Example

Phase 2 (undo)

Database

A	320
B	400
C	100
D	60
E	120

undo list
{}

Log records

[T1 start]
[T1, B, 200, 400]
[T1, C, 50, 100]
[T2 start]
[T2, E, 120, 480]
[T1, A, 100, 320]
[T1 commit]
[T2, A, 320, 520]
[T2, D, 60, 530]
[T2, D, 60]
[T2, A, 320]
[T2, E, 120]
[T2, abort]



Summary

- Recovery: ensuring atomicity and durability despite failures and crashes
- WAL rule
- Log-based recovery
 - All changes need to be written into the log file
 - A transaction commits when the commit entry in the log is written